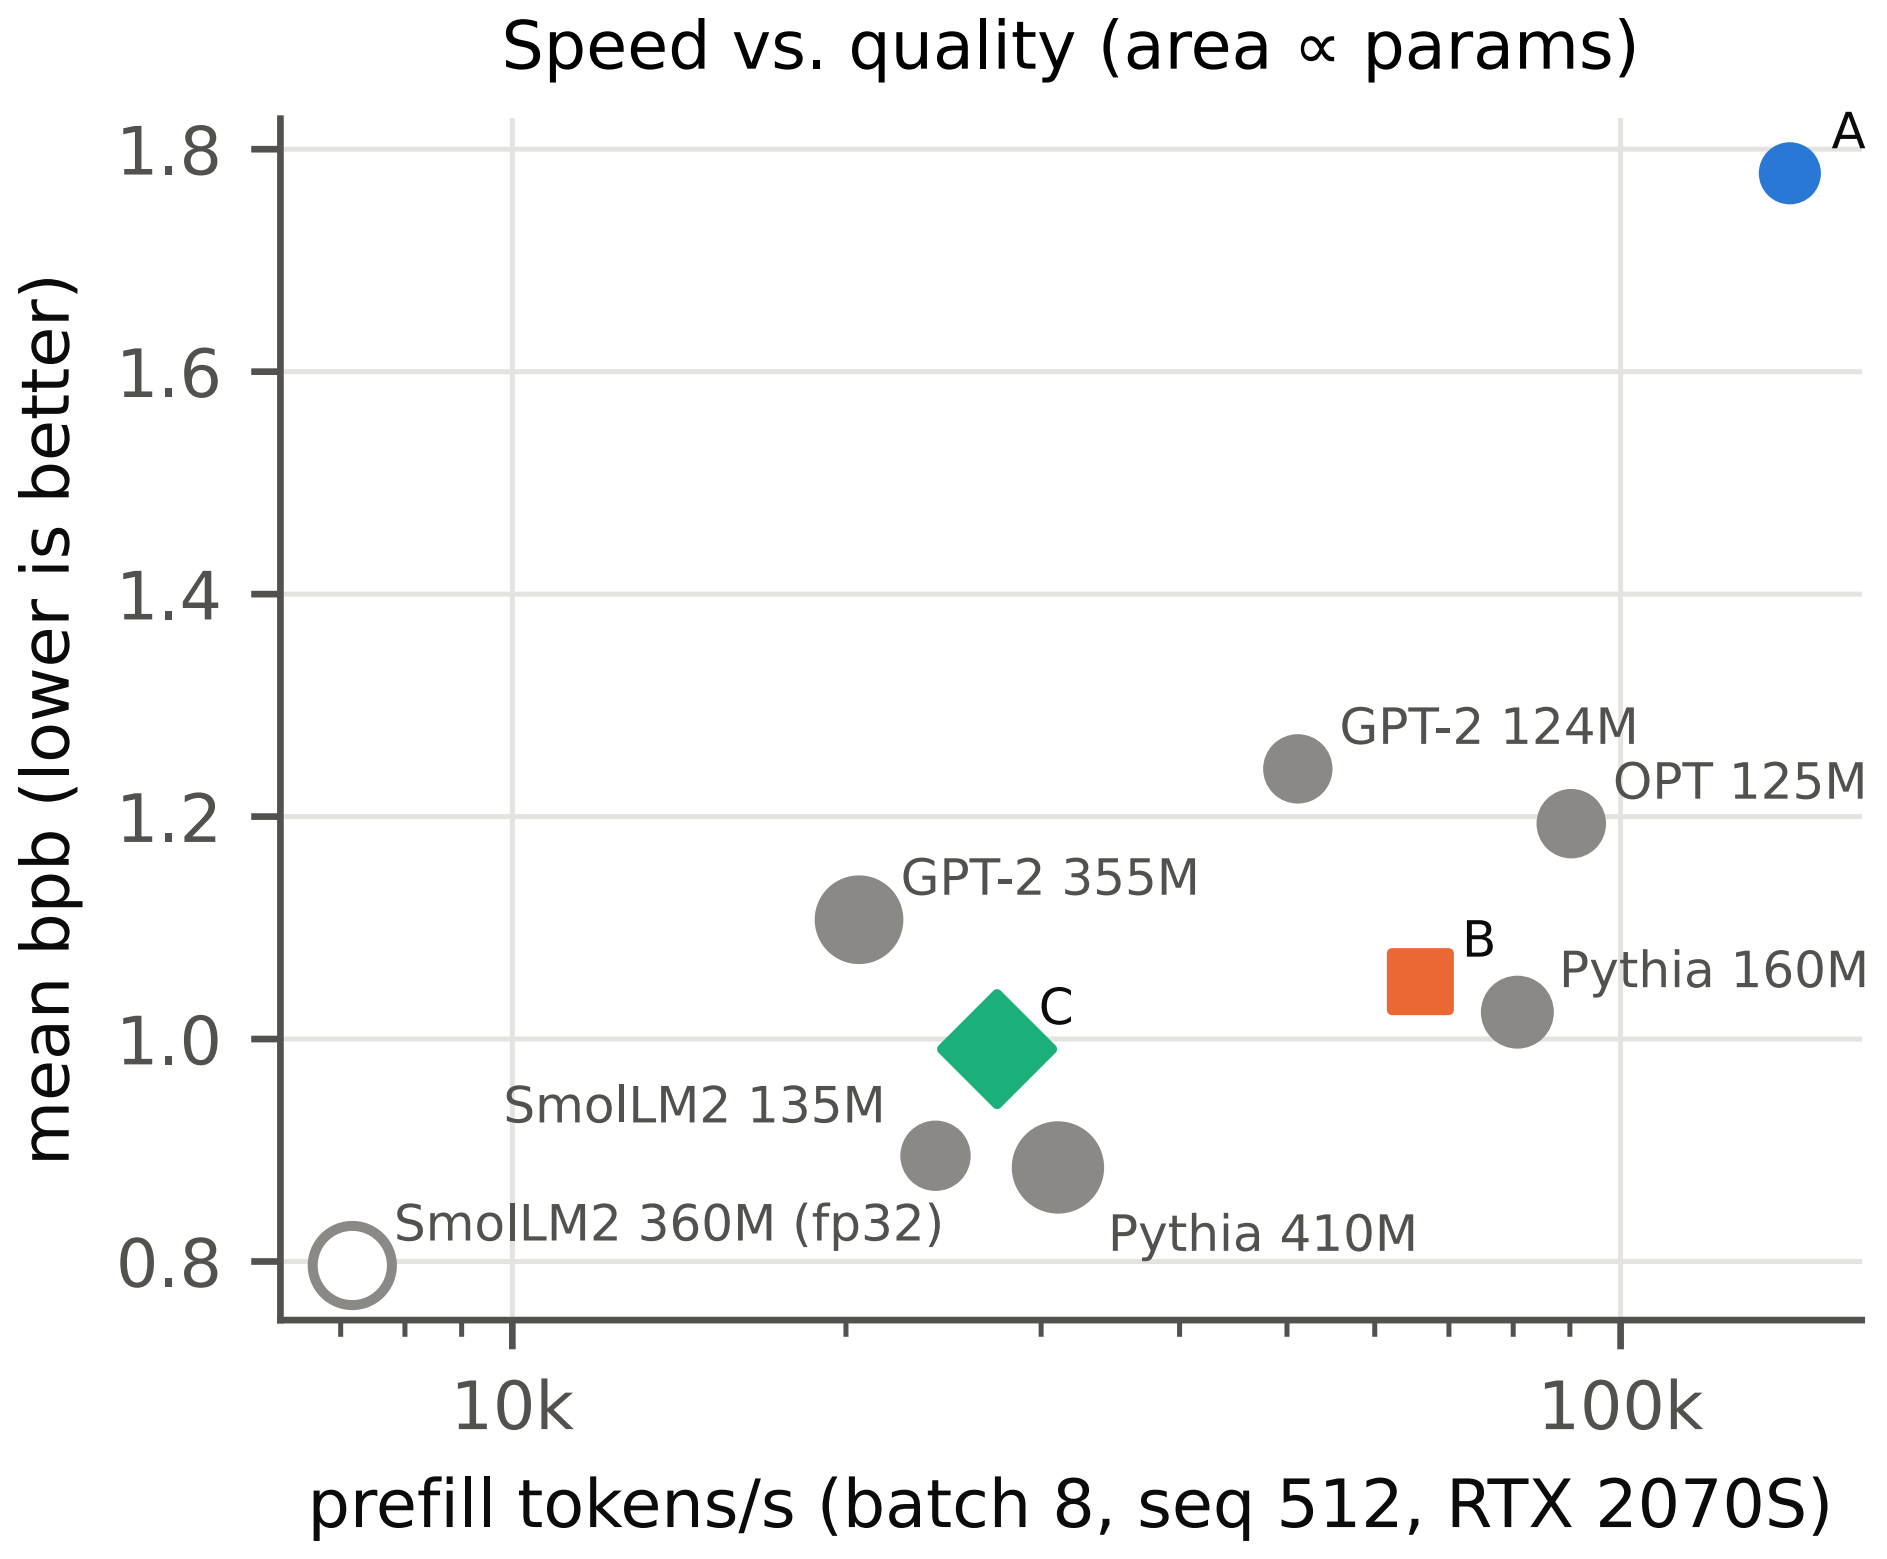
Speed vs. quality (area ∝ params)
mean bpb (lower is better)
1.8
1.6
1.4
1.2
1.0
0.8
A
GPT-2 124M
OPT 125M
GPT-2 355M
B
Pythia 160M
C
SmolLM2 135M
SmolLM2 360M (fp32)
Pythia 410M
10k
100k
prefill tokens/s (batch 8, seq 512, RTX 2070S)