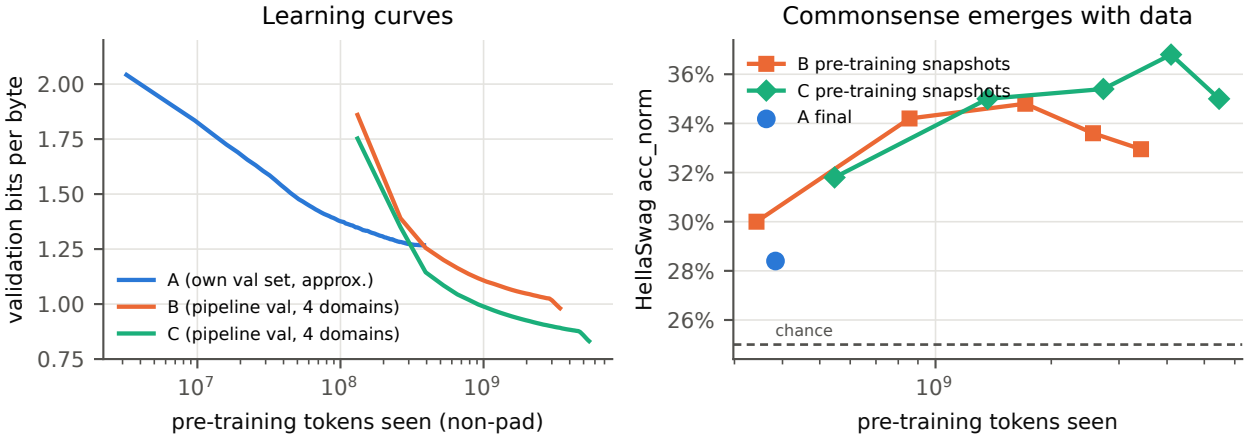
Learning curves
validation bits per byte
2.00
1.75
1.50
1.25
1.00
0.75
A (own val set, approx.)
B (pipeline val, 4 domains)
C (pipeline val, 4 domains)
10^7
10^8
10^9
pre-training tokens seen (non-pad)
Commonsense emerges with data
HellaSwag acc_norm
36%
34%
32%
30%
28%
26%
B pre-training snapshots
C pre-training snapshots
A final
chance
10^9
pre-training tokens seen