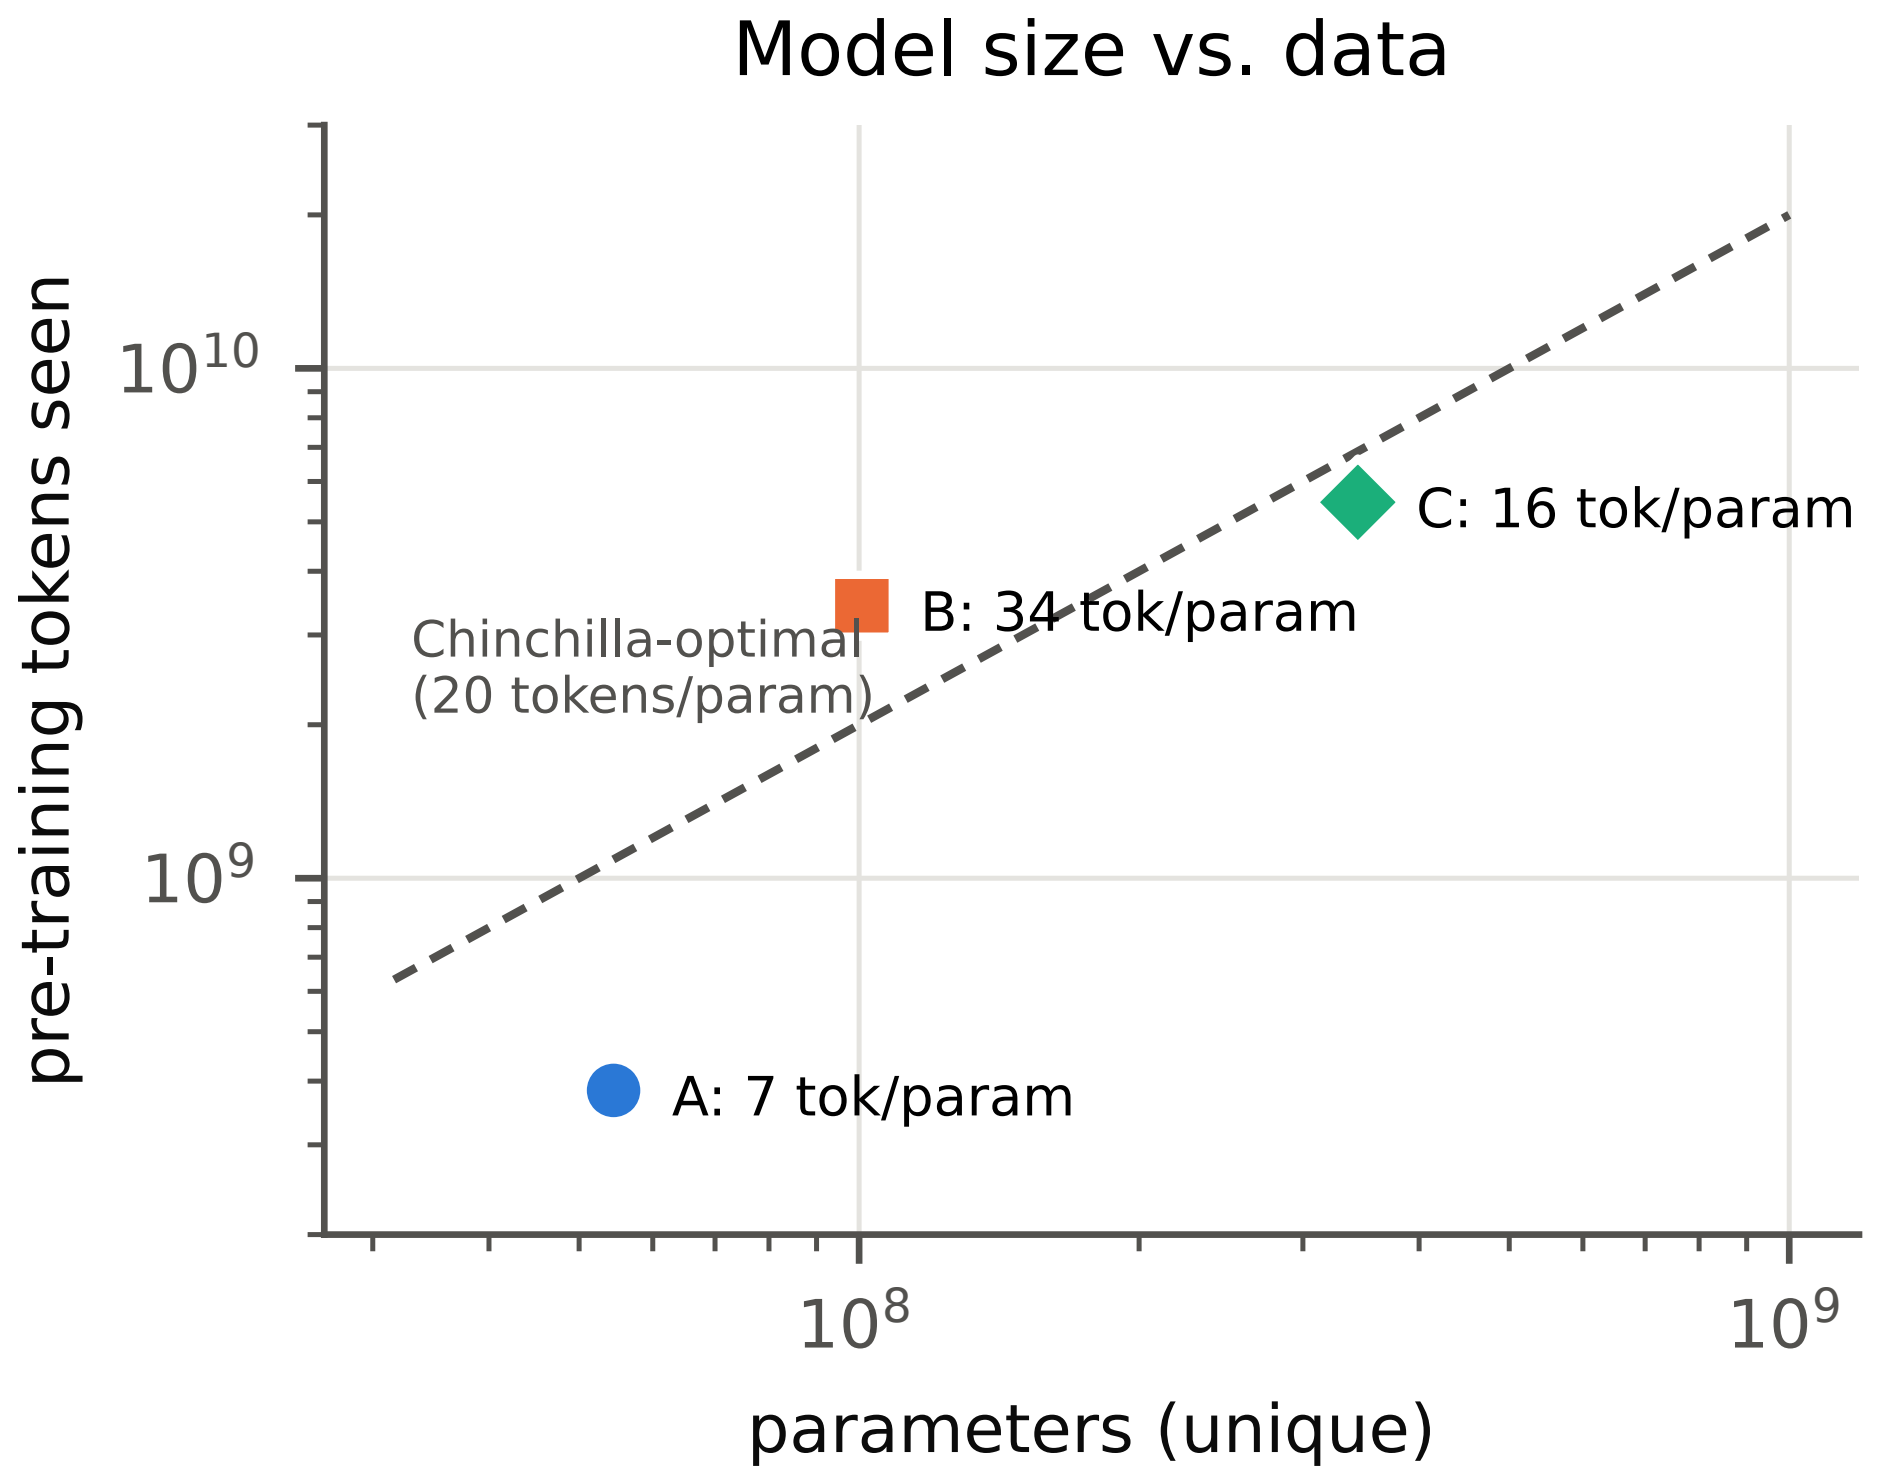
Model size vs. data
pre-training tokens seen
10^{10}
10^{9}
Chinchilla-optimal
(20 tokens/param)
B: 34 tok/param
C: 16 tok/param
A: 7 tok/param
10^{8}
10^{9}
parameters (unique)