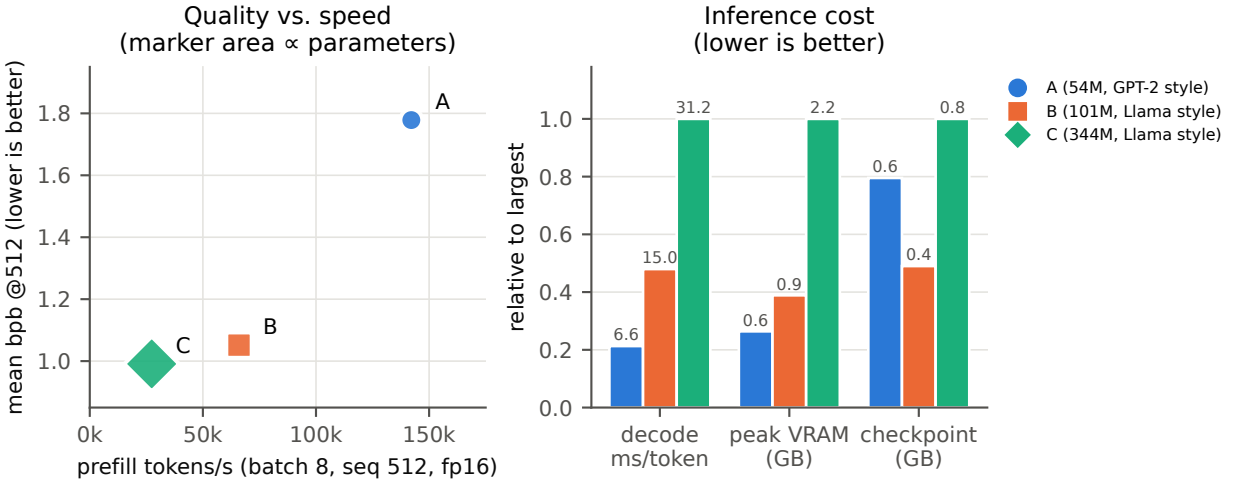
Quality vs. speed
(marker area ∝ parameters)
mean bpb @512 (lower is better)
1.8
1.6
1.4
1.2
1.0
A
B
C
0k
50k
100k
150k
prefill tokens/s (batch 8, seq 512, fp16)
Inference cost
(lower is better)
relative to largest
1.0
0.8
0.6
0.4
0.2
0.0
31.2
2.2
0.8
15.0
0.9
0.4
6.6
0.6
0.6
decode
ms/token
peak VRAM
(GB)
checkpoint
(GB)
A (54M, GPT-2 style)
B (101M, Llama style)
C (344M, Llama style)